

DEEP LEARNING PROJECT

IMPLEMENT A DEEP LEARNING MODEL FOR IMAGE CLASSIFICATION OR NATURAL LANGUAGE PROCESSING USING TENSORFLOW OR PYTORCH

✓ IMPORTING LIBRARIES & LOADING DATASET

```
import tensorflow as tf
from tensorflow.keras import layers, models
import matplotlib.pyplot as plt
import numpy as np
```

```
print("Loading CIFAR-10 dataset...")
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
```

↻ Loading CIFAR-10 dataset...
 Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
 170498071/170498071 ————— 3s 0us/step

✓ NORMALISATION

```
x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',  
               'dog', 'frog', 'horse', 'ship', 'truck']
```

✓ Sample images

```
plt.figure(figsize=(10, 5))
for i in range(10):
    plt.subplot(2, 5, i+1)
    plt.imshow(x_train[i])
    plt.title(class_names[y_train[i][0]])
    plt.axis("off")
plt.show()
```




✓ BUILD CNN MODEL

```
model = models.Sequential([
```

```


layers.Conv2D(32, (3,3), activation='relu', input_shape=(32, 32, 3)),
layers.MaxPooling2D((2,2)),
layers.Conv2D(64, (3,3), activation='relu'),
layers.MaxPooling2D((2,2)),
layers.Conv2D(64, (3,3), activation='relu'),
layers.Flatten(),
layers.Dense(64, activation='relu'),
layers.Dense(10)
])

```

 /usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape` to `input_shape` in the constructor of `Conv2D` or `Conv3D` layers. It is deprecated and will be removed in Keras 3.0.0. Use `input_shape` in the `compile` method instead.

✓ SUMMARY

```
model.summary()
```

 Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
conv2d_2 (Conv2D)	(None, 4, 4, 64)	36,928
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 64)	65,600
dense_1 (Dense)	(None, 10)	650

Total params: 122,570 (478.79 KB)
 Trainable params: 122,570 (478.79 KB)
 Non-trainable params: 0 (0.00 B)

✓ COMPILE MODEL

```

model.compile(optimizer='adam',
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

```

✓ TRAIN MODEL

```

history = model.fit(x_train, y_train, epochs=10,
                    validation_data=(x_test, y_test))

```

 Epoch 1/10
 1563/1563 ————— 67s 42ms/step - accuracy: 0.3393 - loss: 1.7783 - val_accuracy: 0.5500 - val_loss: 1.2486
 Epoch 2/10
 1563/1563 ————— 79s 40ms/step - accuracy: 0.5750 - loss: 1.1929 - val_accuracy: 0.6182 - val_loss: 1.0668
 Epoch 3/10
 1563/1563 ————— 81s 40ms/step - accuracy: 0.6429 - loss: 1.0176 - val_accuracy: 0.6446 - val_loss: 1.0043
 Epoch 4/10
 1563/1563 ————— 61s 39ms/step - accuracy: 0.6804 - loss: 0.9083 - val_accuracy: 0.6844 - val_loss: 0.9207
 Epoch 5/10
 1563/1563 ————— 83s 40ms/step - accuracy: 0.7149 - loss: 0.8213 - val_accuracy: 0.6953 - val_loss: 0.8870
 Epoch 6/10
 1563/1563 ————— 62s 40ms/step - accuracy: 0.7336 - loss: 0.7623 - val_accuracy: 0.6985 - val_loss: 0.8695
 Epoch 7/10
 1563/1563 ————— 81s 39ms/step - accuracy: 0.7499 - loss: 0.7090 - val_accuracy: 0.7029 - val_loss: 0.8542
 Epoch 8/10
 1563/1563 ————— 82s 39ms/step - accuracy: 0.7649 - loss: 0.6699 - val_accuracy: 0.7041 - val_loss: 0.8660
 Epoch 9/10
 1563/1563 ————— 82s 39ms/step - accuracy: 0.7845 - loss: 0.6170 - val_accuracy: 0.7134 - val_loss: 0.8483

Epoch 10/10

1563/1563 — 61s 39ms/step - accuracy: 0.7932 - loss: 0.5867 - val_accuracy: 0.7162 - val_loss: 0.8501

✓ EVALUATE MODEL

```
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print(f'\nTest Accuracy: {test_acc*100:.2f}%')
```

```
313/313 - 3s - 10ms/step - accuracy: 0.7162 - loss: 0.8501
```

Test Accuracy: 71.62%

✓ VISUALIZE TRAINING RESULTS

```
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.show()
```

